StackCube: learning a stronger PPO recipe

Same 24,985,600 training transitions per seed · 3 random initializations · L4 GPU

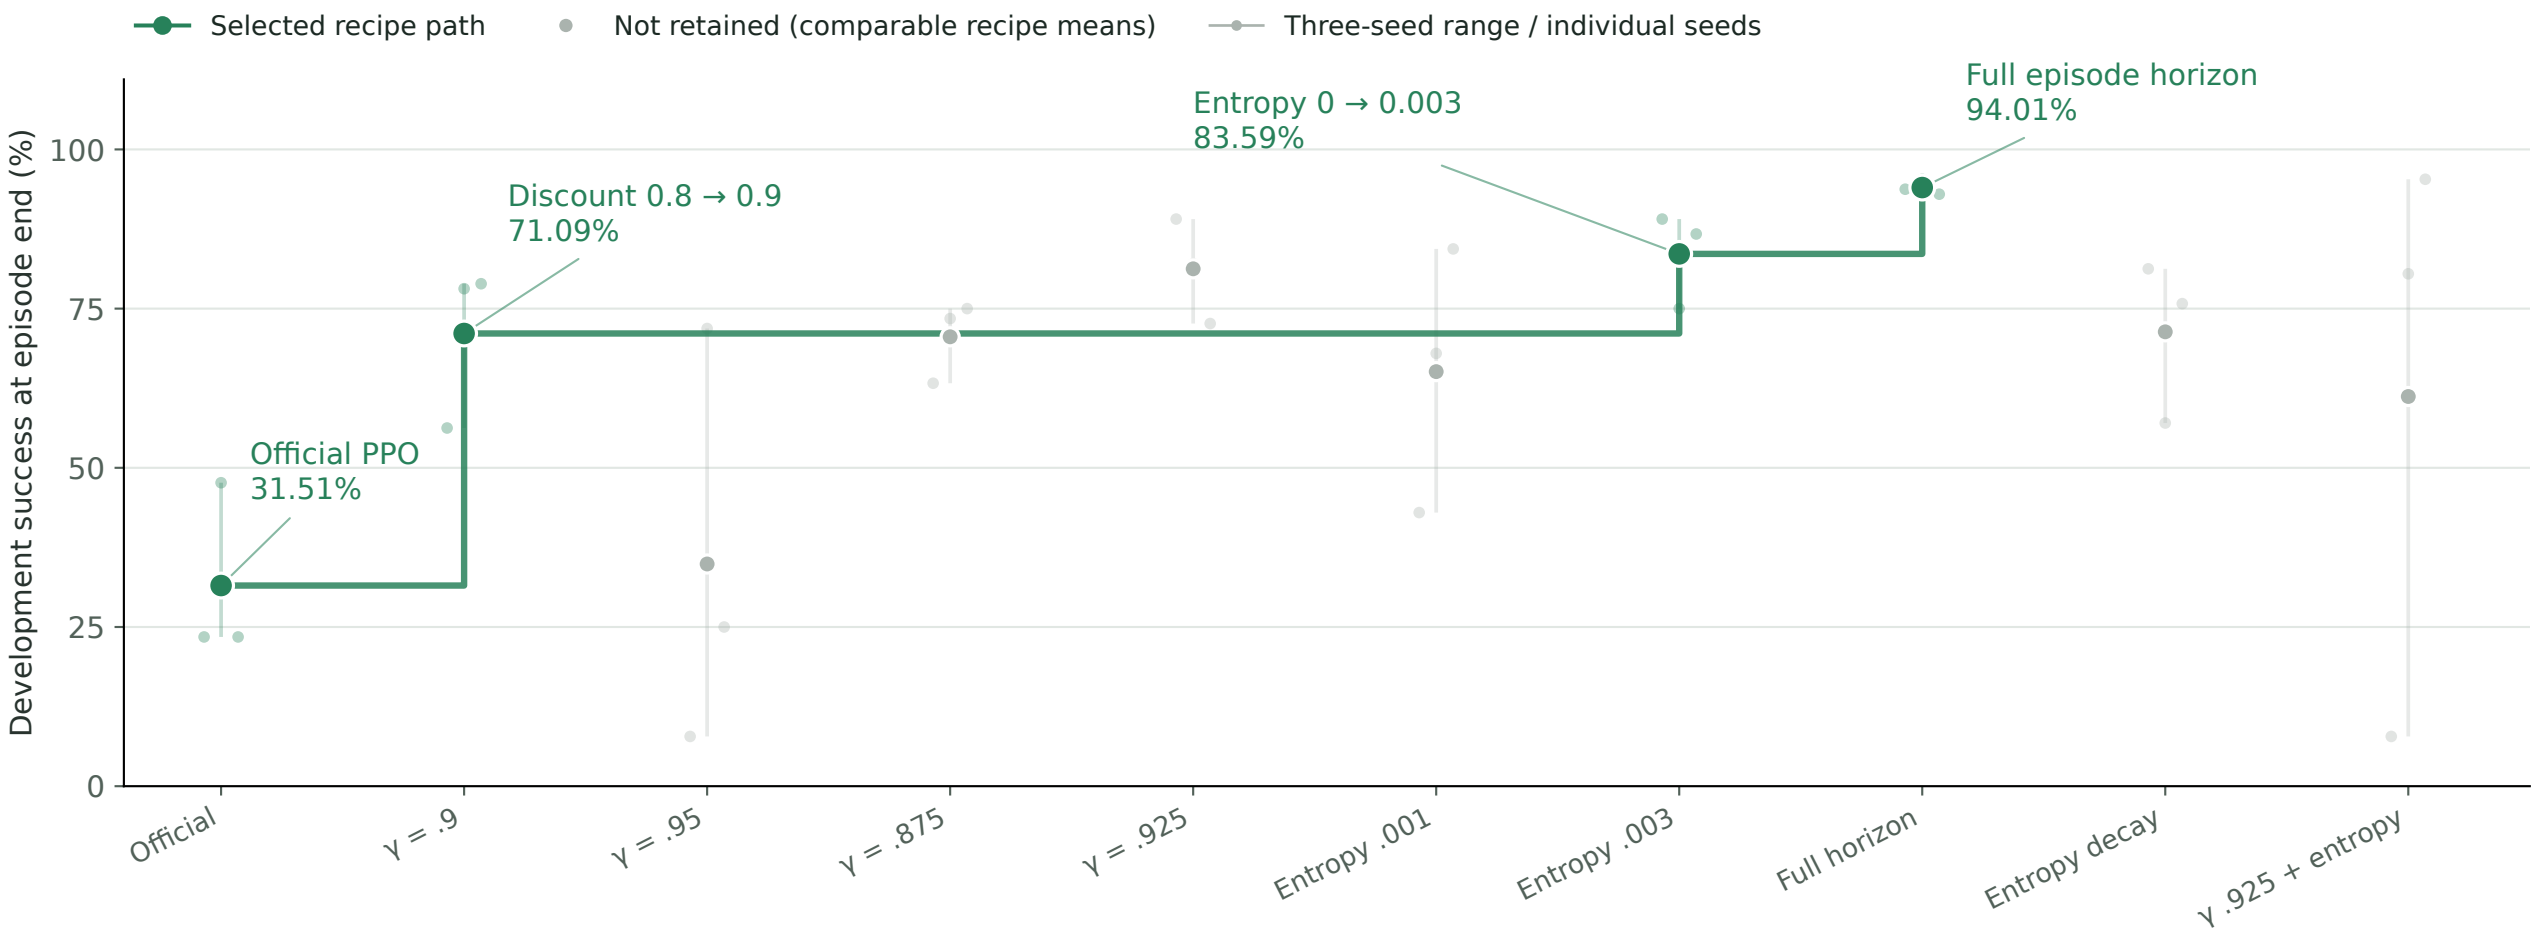

128 fixed development episodes per trained policy. Points are means across 3 training seeds; ranges are not confidence intervals.
Selected path ≠ running best: the gray γ = .925 comparison exceeded the earlier selected recipe. No smoothing or interpolation.

Independent confirmation Official 32.23% → discount + entropy 82.10% → full horizon 94.08%

3 training seeds × 512 confirmation episodes per recipe; discount-only recipe was not confirmed. Full horizon disables partial resets.